

MediaMarkt Hackathon - Solution to the Cloud Challenge

By Mario Valentin Ochoa Mota (@mvochoa)

All the files and instructions for this challenge are in the repository <https://github.com/mvochoa/mediamarkt-hackathon-cloud>.

Configurations.....	1
Configuration of the working environment.....	1
Application configuration.....	2
Google Cloud Configuration.....	2
Application of terraform manifests.....	3
(IMPORTANT) Problem with permissions in Google Cloud.....	5
Solution for the lack of permissions.....	6
Testing the operation of the cloud build.....	7
Docker Compose file.....	9
Questions about roles.....	10
DevOps.....	10
Finances.....	10

Configurations

It is necessary to configure the working environment, the application repository to be CI/CD and the google cloud console.

Configuration of the working environment

It is necessary to make a clone of the repository <https://github.com/mvochoa/mediamarkt-hackathon-cloud> if you open it with the visual studio code editor and have installed the [Dev Containers](#) extension, the editor is configured with everything that is required to work.

In case you do not use the editor and extension, the binaries required are:

- Terraform: => v1.4.2
- Google Cloud CLI (gcloud): The latest version

Application of terraform manifests

All the creation and configuration of Google Cloud resources is done with terraform files, also the Kubernetes manifest (deployments, services, etc) is done with the terraform provider for kubernetes.

First you need to create a bucket to store the terraform `.tfstate` files by executing the following commands:

```
$ cd src/init/
$ terraform init
$ terraform apply
...
Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

Enter a value: yes

...
```

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

Outputs:

```
bucket = "08fe34899c02123a-tfstate"
```

The resources created are:

Name	Type	Description
google_storage_bucket.default	resource	Bucket where the terraform <code>.tfstate</code> files will be stored
random_id.bucket_prefix	resource	It is used to generate a random number for the bucket name

Now once the bucket is created we can apply the rest of the manifests with the commands:

```
$ cd ..
$ terraform init
$ terraform apply
...
Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.
```

Enter a value: yes

...

Apply complete! Resources: 8 added, 0 changed, 0 destroyed.

Outputs:

service = "http://130.211.57.127"

The resources created are:

Name	Type	Description
google_artifact_registry_repository.mediarmarkt-cloud	resource	Repository where the docker images built by the cloud build will be stored
google_cloudbuild_trigger.mms-cloud-skeleton	resource	The Cloud Build that is executed each time the main branch of the repository is pushed
google_container_cluster.mediarmarkt-cloud	resource	
google_container_node_pool.mediarmarkt-cloud	resource	
kubernetes_config_map_v1.continuous-delivery	resource	Kubernetes manifest containing the script to make the image change in the deployment when a new image is uploaded to the artifact registry
kubernetes_cron_job_v1.continuous-delivery	resource	Kubernetes manifest that creates a pod every 5 minutes to validate if a new image exists and change the image on deployment
kubernetes_deployment_v1.mms-cloud-skeleton	resource	Kubernetes manifest for application deployment generates only 2 pods
kubernetes_service_v1.mms-cloud-skeleton	resource	Kubernetes Manifest for the service that will generate a Load Balancer with a Public IP so that the application can be accessed.
google_client_config.default	data source	

google_project.project	data source	
----------------------------------------	-------------	--

(IMPORTANT) Problem with permissions in Google Cloud

Due to lack of permissions for the google cloud IAM with the user assigned to me by the Nuwe team, the cloud build cannot update the deployment image in kubernetes. The resources that were not applied, are commented in the `src/cloudbuild.tf` file.

The resources that cannot be created are the assignment of the kubernetes role to the cloud build service account and the commands to make the image change in the kubernetes deployment:

```
resource "google_project_iam_member" "cloudbuild-kubernetes" {
  project = local.project
  role    = "roles/container.developer"
  member  = "serviceAccount:${data.google_project.project.number}@cloudbuild.gserviceaccount.com"
}

resource "google_cloudbuild_trigger" "mms-cloud-skeleton" {
  ...

  build {
    ...

    step {
      name          = "gcr.io/google.com/cloudsdktool/cloud-sdk"
      entrypoint    = "gcloud"
      args          = ["container", "clusters", "get-credentials", local.cluster_name,
"--region=${local.region}", "--project=${local.project}"]
    }

    step {
      name          = "gcr.io/google.com/cloudsdktool/cloud-sdk"
      entrypoint    = "kubectl"
      args          = ["set", "image", "deployment/${local.repository_name}",
"${local.repository_name}=${local.docker_image_name}"]
    }
  }
}
```

The commands that the cloud build will not be able to execute are:

```
$ gcloud container clusters get-credentials mediamarkt-cloud --region=europe-west1
--project=apmxlt6gss96qvmj7t3s7eexynz6ub
```

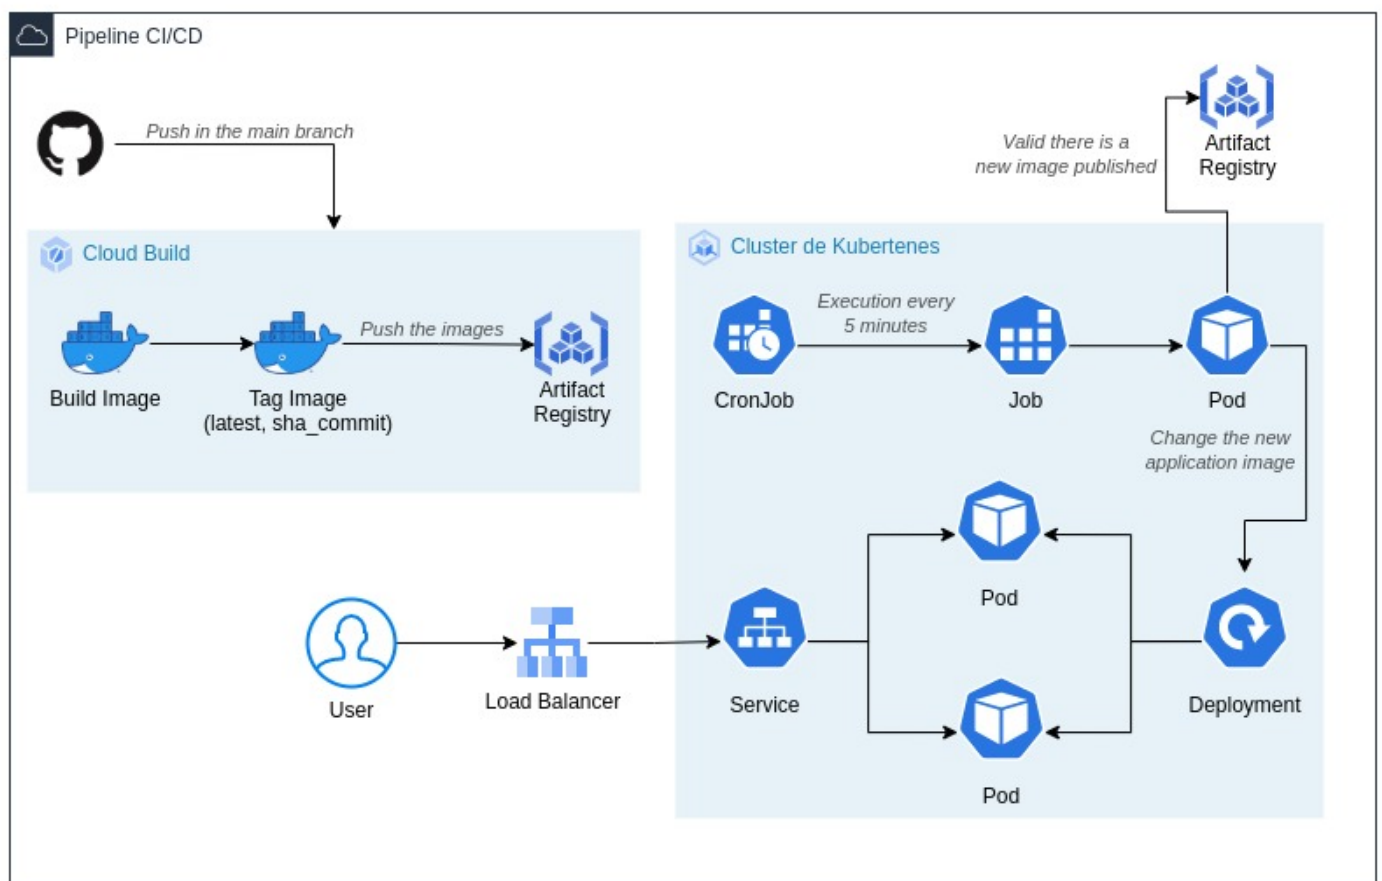
```
$ kubectl set image deployment/mms-cloud-skeleton  
mms-cloud-skeleton=europe-west1-docker.pkg.dev/apmxlt6gss96qvmj7t3s7eexynz6ub/mediamarkt-cloud/mm  
s-cloud-skeleton:$SHORT_SHA
```

Solution for the lack of permissions

I found this solution by reviewing the service account of the kubernetes nodes, I found that it has permission to access the cluster, so internally in the cluster can run the commands that can not be executed in the cloud build for lack of permissions.

The solution was to create a cronjob in kubernetes that checks every 5 minutes, if a new image was uploaded to the artifact registry and when it detects a new image it updates the deployment with that new image. The script and the kubernetes manifests are in [src/continuous_delivery.kube.tf](https://github.com/mediamarkt-cloud/mms-cloud-skeleton/blob/main/src/continuous_delivery.kube.tf)

Applying this solution the CI/CD diagram looks like this:



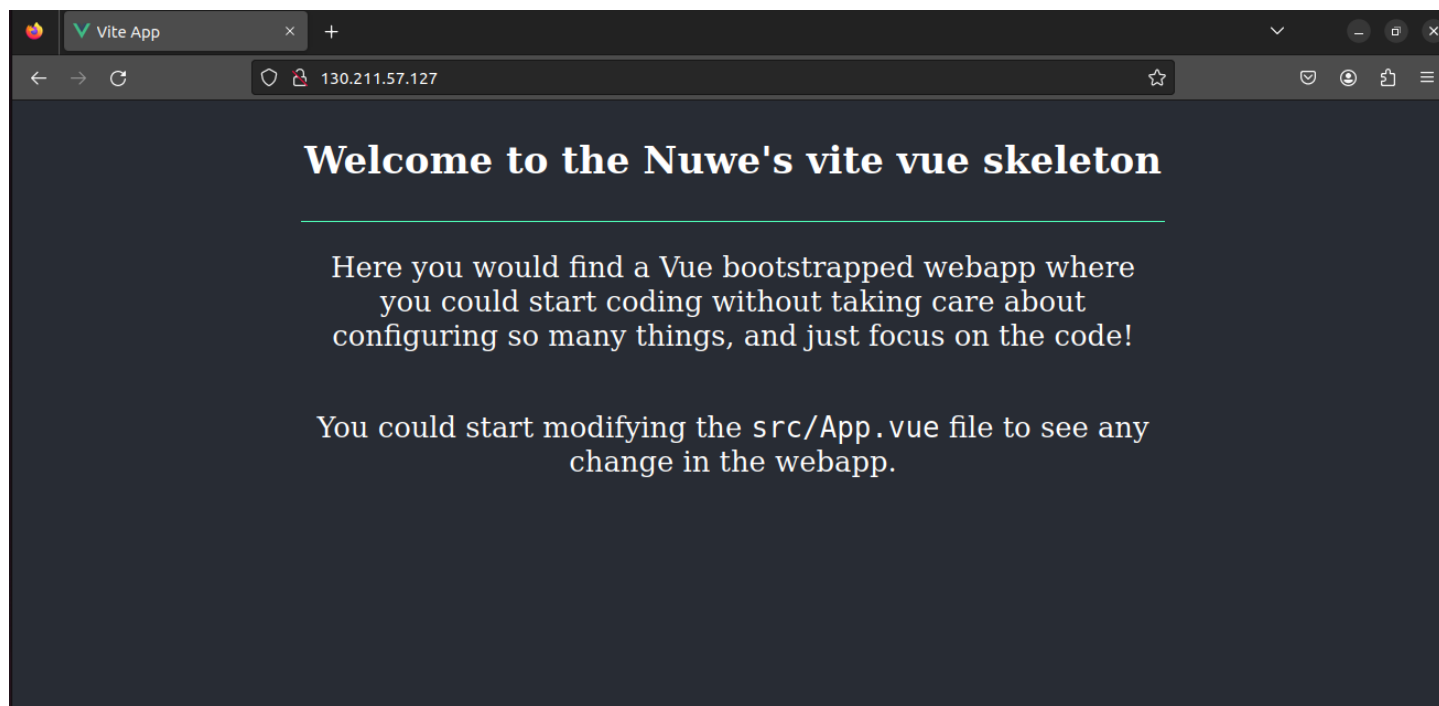
This way the CI/CD is complete 🎉🎉🎉. I know it's not a very clean way to solve the problem but since I couldn't assign permissions correctly, it was the only way I found to get the CI/CD to work.

Testing the operation of the cloud build

A manual execution of the cloud build is going to be done to see the application as it is at the moment.

The first screenshot shows the 'Ejecuta activadores' (Run triggers) dialog in the Google Cloud console. It is configured for the 'mms-cloud-skeleton' repository, with the 'main' branch and a 'Confirmar hash' option selected. The second screenshot shows the 'Detalles de compilación' (Build details) page for build '448a2b05'. It lists four steps: 1. 'gcr.io/cloud-builders/docker build -f europe-west1-docker...', 2. 'gcr.io/cloud-builders/docker push europe-west1-docker.p...', 3. 'gcr.io/cloud-builders/docker push europe-west1-docker.p...', and 4. 'gcr.io/cloud-builders/docker push europe-west1-docker.p...'. The third screenshot shows the 'Resúmenes de mms-cloud-skeleton' page in Artifact Registry. It displays a table with one entry: '03b78e3c39cb' with a hash of '842629a' and a 'latest' tag. A notification banner at the top encourages activating vulnerability analysis.

A few minutes later, the web can be viewed on the IP 130.211.57.127 of the kubernetes service.



Now we test uploading changes to the repository to see if the whole pipeline works correctly, the change that was added was:

```
src/ui/Home/HomeComponent.vue
-11,6 +11,7
11 11      You could start modifying the <code>src/App.vue</code> file to see any
12 12      change in the webapp.
13 13  </p>
14 14  + <small>CI/CD created by @mvochoa</small>
14 15  </div>
15 16  </template>
16 17
```

Google Cloud Console - Detalles de compilación

Correcta: 1197/ca3

Activador: mms-cloud-skeleton

Fuente: mvochoa/mms-cloud-skeleton

Rama: main

Confirmar: 672493610

Paso	Duración	REGISTRO DE LA COMPILACIÓN	DETALLES DE LA EJECUCIÓN	ARTIFACTOS DE COMPILACIÓN
Resumen de la compilación				
4 pasos				
0: gcr.io/cloud-builders/docker...	00:01:38			
1: gcr.io/cloud-builders/docker...	00:00:00			
2: gcr.io/cloud-builders/docker...	00:00:02			
3: gcr.io/cloud-builders/docker...	00:00:00			

Google Cloud Console - Resúmenes de mms-cloud-skeleton

Artifacts Registry

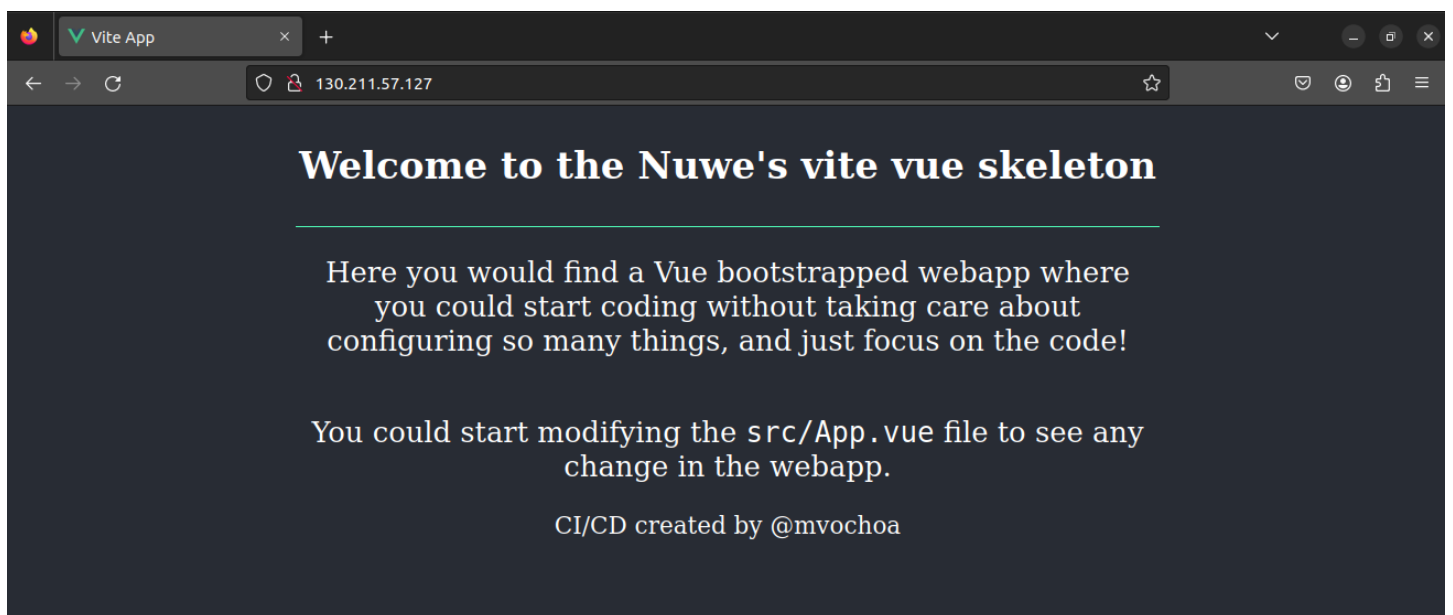
Resúmenes de mms-cloud-skeleton

Configuración

Activar el análisis de vulnerabilidades

Activar

Nombre	Descripción	Etiquetas	Fecha de creación	Actualizado
35c97148a6b5		672493610	latest	Hace unos instantes
55b67bdc39cb		8426294		Hace 4 horas



Docker Compose file

To the github repository <https://github.com/mvochoa/mms-cloud-skeleton> the docker-compose.yaml file was added and the Dockerfile was improved.

The docker-compose.yaml file was optimized so that developers can work faster and more comfortably since a development server is running and the node_modules are stored in volume.

version: '3.8'


```

services:
  app:
    build:
      context: .
      target: builder
    ports:
      - 3000:3000
    volumes:
      - ./app
      - node_modules:/app/node_modules
    command: sh -c 'yarn install && yarn run dev'

```

```

volumes:
  node_modules:

```

The Dockerfile was also adjusted to make the final image lighter and with better performance.

```

...  ...  @@ -1,11 +1,11 @@
1  - FROM node:alpine
2  -
3  - ADD package.json /app/package.json
4  + FROM node:alpine AS builder
5  2
6  3  WORKDIR /app
7  4
8  5  RUN yarn add vite
9  6
10 7  + ADD package.json /app/package.json
11 8  +
12 9  # Installing packages
13 10 RUN yarn install
14 11
15  ...  @@ -16,4 +16,15 @@ RUN yarn build-only
16 16
17 17  EXPOSE 3000
18 18
19  - CMD [ "yarn", "preview" ]
20  + CMD [ "yarn", "dev" ]
21  +
22  + FROM nginx:stable-alpine
23  +
24  + ENV NGINX_PORT=3000
25  +
26  + COPY --from=builder /app/dist /usr/share/nginx/html
27  +
28  + RUN mkdir /etc/nginx/templates/ \
29  +   && sed 's/listen 80/listen ${NGINX_PORT}g' /etc/nginx/conf.d/default.conf >
30  +   /etc/nginx/templates/default.conf.template
31  +
32  + EXPOSE 3000

```

Questions about roles

DevOps

For the management of the kubernetes cluster in Google Cloud you have to assign to the group, the role of **Kubernetes Engine Developer** this role allows the management of the different services that the cluster has.

Finances

To manage the billing in Google Cloud, the **Billing Account Viewer** role must be assigned to the group, this role only allows the management of billing.